

Policy Gradient Methods

wonseok Jung

1. Two methods of choosing action

- action-value :
 - Learning the action value
 - Estimate action value을 바탕으로 action을 선택한다.
 - Policies would not even exist without the action-value estimates
- Parameterized policy :
 - select actions without consulting value function
 - Value function still be used to learn policy parameter
 - Value function이 action을 선택하는 기준으로 사용되지 않는다

1.1 Policy parameter vector

- $\theta \in R^{d'}$:Policy's parameter vector
- $\pi(a \mid s, \theta) = Pr \{A_t = a \mid S_t = S, \theta_t = \theta\}$
 - Probability action a is taken at time t , given that the environment is in state s at time t with parameter θ

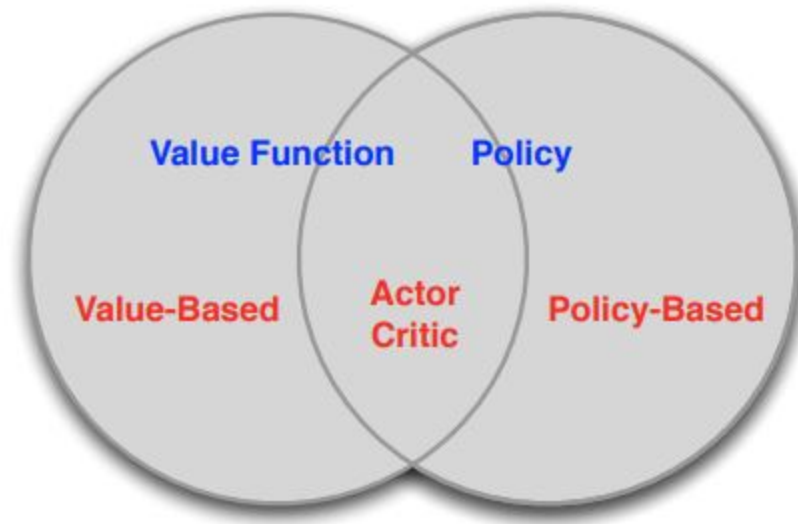
1.2 Learning policy parameter

$$\theta_{t+1} = \theta_t + \alpha \widehat{\nabla J(\theta)}$$

- $\nabla J(\theta_t)$: stochastic estimate
 - expectation approximates the gradient of the measure to its argument θ

- Policy gradient methods : optimizing parametrized policies with respect to the expected return by gradient ascent
- Actor-critic methods : Methods that learn approximations to both policy and value function

- Value Based
 - Learnt Value Function
 - Implicit policy (e.g. ϵ -greedy)
- Policy Based
 - No Value Function
 - Learnt Policy
- Actor-Critic
 - Learnt Value Function
 - Learnt Policy



2. Policy Approximation and ITS Advantages

3. The Policy gradient Theorem

- $\epsilon - greedy$ 는 작은 action value의 변화에도 action 선택이 완전히 바뀔수 있다.
- 하지만 policy parameterization 방법은 Parameter를 배우며 Policy parameter의 action probabilities가 smoothly하게 변한다.
- 이렇게 parameter에 의해 policy가 달라진다면, policy-gradient 방법을 이용하여 gradient ascent를 approximate하는것이 가능하다.

3.1 Two type of cases

- $J(\theta)$: Performance measure
- Episodic case, Continuous case 두 가지로 나눌수 있다.
- Episodic case : the performance measure as the value of the start state of the episode,

$$J(\theta) \doteq v_{\pi_\theta}(s_0)$$

- $v_{\pi_\theta}(s_0)$: True value for π_θ , the policy determined by θ

3.2 Challenging

- With function approximation, it may seem challenging to change the policy parameter in a way that ensures improvements.
 - performance depends on both the action selection and the distribution of states in which those selections are made
 - Both of these are affected by the policy parameter
- Policy gradient theorem : analytic expression for the gradient of performance with respect to the policy parameter

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a \mid s, \theta)$$

3.3 Policy Gradient

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a \mid s, \theta)$$

π : policy corresponding to parameter vector θ

\propto : propotional to

μ : on-policy distribuion under

- From chapter 10 :

- $\mu_\pi(s) \doteq \lim_{t \rightarrow \infty} Pr \{S_t = s \mid A_{0:t-1} \sim \pi\}$
- steady-state distribution

4. REINFORCE: Monte Carlo Policy Gradient

$$\begin{aligned}\nabla J(\theta) &\propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a \mid s, \theta) \\ &= E_\pi \left[\sum_a q_\pi(S_t, a) \nabla \pi(a \mid S_t, \theta) \right]\end{aligned}$$

- Policy gradient theorem : sum over a states weight by how often the states occur under the target policy π

4.1 Replacing a with the sample action A_t

•

$$\begin{aligned}\nabla J(\boldsymbol{\theta}) &= \mathbb{E}_{\pi} \left[\sum_a \pi(a|S_t, \boldsymbol{\theta}) q_{\pi}(S_t, a) \frac{\nabla \pi(a|S_t, \boldsymbol{\theta})}{\pi(a|S_t, \boldsymbol{\theta})} \right] \\ &= \mathbb{E}_{\pi} \left[q_{\pi}(S_t, A_t) \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta})}{\pi(A_t|S_t, \boldsymbol{\theta})} \right] && \text{(replacing } a \text{ by the sample } A_t \sim \pi) \\ &= \mathbb{E}_{\pi} \left[G_t \frac{\nabla \pi(A_t|S_t, \boldsymbol{\theta})}{\pi(A_t|S_t, \boldsymbol{\theta})} \right], && \text{(because } \mathbb{E}_{\pi}[G_t|S_t, A_t] = q_{\pi}(S_t, A_t))\end{aligned}$$

- G_t : Return

4.2 REINFORCE algorithm

- Updte Rule :

$$\theta_{t+1} \doteq \theta_t + \alpha G_t \frac{\nabla \pi(A_t | S_t, \theta_t)}{\pi(A_t | S_t, \theta_t)}$$

- REINFORCE uses the comple te return from time t.
- All future rewards update until the end of the epsiode.
- REINFORCE는 MonteCalo 알고리즘을 사용한다. 모든 업데이트는 episode가 끝난뒤 이루어진다.

4.3 REINFORCE-Monte Carlo pseudocode

REINFORCE: Monte-Carlo Policy-Gradient Control (episodic) for π_*

Input: a differentiable policy parameterization $\pi(a|s, \theta)$

Algorithm parameter: step size $\alpha > 0$

Initialize policy parameter $\theta \in \mathbb{R}^{d'}$ (e.g., to $\mathbf{0}$)

Loop forever (for each episode):

 Generate an episode $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$, following $\pi(\cdot|\cdot, \theta)$

 Loop for each step of the episode $t = 0, 1, \dots, T - 1$:

$$G \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k \quad (G_t)$$

$$\theta \leftarrow \theta + \alpha \gamma^t G \nabla \ln \pi(A_t | S_t, \theta)$$




Figure 13.1: REINFORCE on the short-corridor gridworld (Example 13.1). With a good step size, the total reward per episode approaches the optimal value of the start state.

- α : step size에 따른 Total reward의 차이 비교

5. REINFORCE with Baseline

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a q_\pi(s, a) \nabla \pi(a \mid s, \theta)$$

- 위의 policy gradient theorem 은 다음의 식과 같이 어떠한 baseline $b(s)$ 에 의해 action value의 비교 식으로 바꿀수 있다.

$$\nabla J(\theta) \propto \sum_s \mu(s) \sum_a (q_\pi(s, a) - b(s)) \nabla \pi(a \mid s, \theta)$$

5.1 REINFORCE update rule with baseline

- $\theta_{t+1} \doteq \theta_t + \alpha(G_t - b(S_t)) \frac{\nabla \pi(A_t|S_t, \theta_t)}{\pi(A_t|S_t, \theta_t)}$
- Baseline을 사용하므로 variance를 줄일수 있다.
- 그러므로 learning 속도가 빨라진다.
- For MDP, baseline should vary with state
 - some states all actions have high values
 - some states all actions have low values

5.3 Baseline




Figure 13.2: Adding a baseline to REINFORCE can make it learn much faster, as illustrated here on the short-corridor gridworld (Example 13.1). The step size used here for plain REINFORCE is that at which it performs best (to the nearest power of two; see Figure 13.1). Each line is an average over 100 independent runs.

5.2 Reinforcement with baseline pseudocode

REINFORCE with Baseline (episodic), for estimating $\pi_{\theta} \approx \pi_*$

Input: a differentiable policy parameterization $\pi(a|s, \theta)$

Input: a differentiable state-value function parameterization $\hat{v}(s, \mathbf{w})$

Algorithm parameters: step sizes $\alpha^{\theta} > 0$, $\alpha^{\mathbf{w}} > 0$

Initialize policy parameter $\theta \in \mathbb{R}^{d'}$ and state-value weights $\mathbf{w} \in \mathbb{R}^d$ (e.g., to $\mathbf{0}$)

Loop forever (for each episode):

 Generate an episode $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$, following $\pi(\cdot|\cdot, \theta)$

 Loop for each step of the episode $t = 0, 1, \dots, T - 1$:

$$G \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k \quad (G_t)$$

$$\delta \leftarrow G - \hat{v}(S_t, \mathbf{w})$$

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha^{\mathbf{w}} \gamma^t \delta \nabla \hat{v}(S_t, \mathbf{w})$$

$$\theta \leftarrow \theta + \alpha^{\theta} \gamma^t \delta \nabla \ln \pi(A_t|S_t, \theta)$$